



ETC3550/ETC5550

Applied forecasting

Ch2. Time series graphics

OTexts.org/fpp3/

Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
```

```
## # Key:           Country [263]
```

##	Year	Country	GDP	Imports	Exports	Population
##	<dbl>	<fct>	<dbl>	<dbl>	<dbl>	<dbl>
## 1	1960	Afghanistan	537777811.	7.02	4.13	8996351
## 2	1961	Afghanistan	548888896.	8.10	4.45	9166764
## 3	1962	Afghanistan	546666678.	9.35	4.88	9345868
## 4	1963	Afghanistan	751111191.	16.9	9.17	9533954
## 5	1964	Afghanistan	800000044.	18.1	8.89	9731361
## 6	1965	Afghanistan	1006666638.	21.4	11.3	9938414
## 7	1966	Afghanistan	1399999967.	18.6	8.57	10152331
## 8	1967	Afghanistan	1673333418.	14.2	6.77	10372630
## 9	1968	Afghanistan	1373333367.	15.2	8.90	10604346
## 10	1969	Afghanistan	1408888922.	15.0	10.1	10854428

```
## # ... with 15,140 more rows
```

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
```

```
## # Key:           Country [263]
```

```
##   Year Country      GDP Imports Exports Population
##   Index <fct>      <dbl>   <dbl>   <dbl>         <dbl>
## 1  1960 Afghanistan 537777811.    7.02    4.13    8996351
## 2  1961 Afghanistan 548888896.    8.10    4.45    9166764
## 3  1962 Afghanistan 546666678.    9.35    4.88    9345868
## 4  1963 Afghanistan 751111191.   16.9    9.17    9533954
## 5  1964 Afghanistan 800000044.   18.1    8.89    9731361
## 6  1965 Afghanistan 1006666638.   21.4   11.3    9938414
## 7  1966 Afghanistan 1399999967.   18.6    8.57   10152331
## 8  1967 Afghanistan 1673333418.   14.2    6.77   10372630
## 9  1968 Afghanistan 1373333367.   15.2    8.90   10604346
## 10 1969 Afghanistan 1408888922.   15.0   10.1   10854428
## # ... with 15,140 more rows
```

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
```

```
## # Key:           Country [263]
```

```
##      Year Country      GDP Imports Exports Population
##      Index  Key      <dbl>   <dbl>   <dbl>         <dbl>
##  1  1960 Afghanistan 5377777811.    7.02    4.13    8996351
##  2  1961 Afghanistan 5488888896.    8.10    4.45    9166764
##  3  1962 Afghanistan 5466666678.    9.35    4.88    9345868
##  4  1963 Afghanistan 7511111191.   16.9    9.17    9533954
##  5  1964 Afghanistan 8000000044.   18.1    8.89    9731361
##  6  1965 Afghanistan 10066666638.   21.4   11.3    9938414
##  7  1966 Afghanistan 13999999967.   18.6    8.57   10152331
##  8  1967 Afghanistan 16733333418.   14.2    6.77   10372630
##  9  1968 Afghanistan 13733333367.   15.2    8.90   10604346
## 10  1969 Afghanistan 14088888922.   15.0   10.1   10854428
## # ... with 15,140 more rows
```

tsibble objects

```
global_economy
```

```
## # A tsibble: 15,150 x 6 [1Y]
```

```
## # Key:          Country [263]
```

```
##      Year Country      GDP Imports Exports Population
```

```
##      Index  Key      Measured variables
```

```
## 1  1960 Afghanistan 537777811.      7.02      4.13      8996351
```

```
## 2  1961 Afghanistan 548888896.      8.10      4.45      9166764
```

```
## 3  1962 Afghanistan 546666678.      9.35      4.88      9345868
```

```
## 4  1963 Afghanistan 751111191.     16.9      9.17      9533954
```

```
## 5  1964 Afghanistan 800000044.     18.1      8.89      9731361
```

```
## 6  1965 Afghanistan 1006666638.     21.4     11.3     9938414
```

```
## 7  1966 Afghanistan 1399999967.     18.6      8.57     10152331
```

```
## 8  1967 Afghanistan 1673333418.     14.2      6.77     10372630
```

```
## 9  1968 Afghanistan 1373333367.     15.2      8.90     10604346
```

```
## 10 1969 Afghanistan 1408888922.     15.0     10.1     10854428
```

```
## # ... with 15,140 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:           Region, State, Purpose [304]
##   Quarter Region   State Purpose   Trips
##   <qtr> <chr>      <chr> <chr>    <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```


tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:           Region, State, Purpose [304]
##   Quarter Region   State Purpose   Trips
##   Index   <chr>    <chr> <chr>    <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:      Region, State, Purpose [304]
##   Quarter Region  State Purpose  Trips
##   Index      Keys      <dbl>
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
## # Key:      Region, State, Purpose [304]
```

```
##   Quarter Region State Purpose Trips
##   Index      Keys          Measure
## 1 1998 Q1 Adelaide SA      Business 135.
## 2 1998 Q2 Adelaide SA      Business 110.
## 3 1998 Q3 Adelaide SA      Business 166.
## 4 1998 Q4 Adelaide SA      Business 127.
## 5 1999 Q1 Adelaide SA      Business 137.
## 6 1999 Q2 Adelaide SA      Business 200.
## 7 1999 Q3 Adelaide SA      Business 169.
## 8 1999 Q4 Adelaide SA      Business 134.
## 9 2000 Q1 Adelaide SA      Business 154.
## 10 2000 Q2 Adelaide SA      Business 169.
## # ... with 24,310 more rows
```

tsibble objects

```
tourism
```

```
## # A tsibble: 24,320 x 5 [1Q]
```

```
## # Key:           Region, State, Purpose [304]
```

```
##   Quarter Region State Purpose Trips
```

```
##   Index      Keys      Measure
```

```
## 1 1998 Q1 Adelaide SA      Business 135.
```

```
## 2 1998 Q2 Adelaide SA      Business 110.
```

```
## 3 1998 Q3 Adelaide SA      Business 166.
```

```
## 4 1998 Q4 Adelaide SA      Business 127.
```

```
## 5 1999 Q1 Adelaide SA      Business 137.
```

```
## 6 1999 Q2 Adelaide SA      Business 200.
```

```
## 7 1999 Q3 Adelaide SA      Business 169.
```

```
## 8 1999 Q4 Adelaide SA      Business 134.
```

```
## 9 2000 Q1 Adelaide SA      Business 154.
```

```
## 10 2000 Q2 Adelaide SA      Business 169.
```

```
## # ... with 24,310 more rows
```

Domestic visitor
nights in thousands
by state/region and
purpose.

tsibble objects

- A `tsibble` allows storage and manipulation of multiple time series in R.
- It contains:
 - ▶ An index: time information about the observation
 - ▶ Measured variable(s): numbers of interest
 - ▶ Key variable(s): optional unique identifiers for each series
- It works with tidyverse functions.

The tsibble index

Example

```
mydata <- tsibble(  
  year = 2012:2016,  
  y = c(123, 39, 78, 52, 110),  
  index = year  
)  
mydata
```

```
## # A tsibble: 5 x 2 [1Y]  
##   year      y  
##   <int> <dbl>  
## 1  2012   123  
## 2  2013    39  
## 3  2014    78  
## 4  2015    52  
## 5  2016   110
```

The tsibble index

Example

```
mydata <- tibble(  
  year = 2012:2016,  
  y = c(123, 39, 78, 52, 110)  
) %>%  
  as_tsibble(index = year)  
mydata
```

```
## # A tsibble: 5 x 2 [1Y]  
##   year      y  
##   <int> <dbl>  
## 1  2012   123  
## 2  2013    39  
## 3  2014    78  
## 4  2015    52  
## 5  2016   110
```

The tsibble index

For observations more frequent than once per year, we need to use a time class function on the index.

```
z
```

```
## # A tibble: 5 x 2
##   Month      Observation
##   <chr>          <dbl>
## 1 2019 Jan           50
## 2 2019 Feb           23
## 3 2019 Mar           34
## 4 2019 Apr           30
## 5 2019 May           25
```


The tsibble index

For observations more frequent than once per year, we need to use a time class function on the index.

```
z %>%  
  mutate(Month = yearmonth(Month)) %>%  
  as_tsibble(index = Month)
```

```
## # A tsibble: 5 x 2 [1M]  
##       Month Observation  
##       <mth>         <dbl>  
## 1 2019 Jan           50  
## 2 2019 Feb           23  
## 3 2019 Mar           34  
## 4 2019 Apr           30  
## 5 2019 May           25
```

The tsibble index

Common time index variables can be created with these functions:

Frequency	Function
Annual	<code>start:end</code>
Quarterly	<code>yearquarter()</code>
Monthly	<code>yearmonth()</code>
Weekly	<code>yearweek()</code>
Daily	<code>as_date()</code> , <code>ymd()</code>
Sub-daily	<code>as_datetime()</code>

Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

Australian prison population



Read a csv file and convert to a tibble

```
prison <- readr::read_csv("data/prison_population.csv")
```

```
## # A tibble: 3,072 x 6
```

```
##   date      state gender legal      indigenous count
##   <date>    <chr> <chr>  <chr>    <chr>         <dbl>
## 1 2005-03-01 ACT    Female Remanded ATSI           0
## 2 2005-03-01 ACT    Female Remanded Other         2
## 3 2005-03-01 ACT    Female Sentenced ATSI           0
## 4 2005-03-01 ACT    Female Sentenced Other         0
## 5 2005-03-01 ACT    Male   Remanded ATSI           7
## 6 2005-03-01 ACT    Male   Remanded Other        58
## 7 2005-03-01 ACT    Male   Sentenced ATSI           0
## 8 2005-03-01 ACT    Male   Sentenced Other         0
## 9 2005-03-01 NSW    Female Remanded ATSI          51
## 10 2005-03-01 NSW    Female Remanded Other       131
## # ... with 3,062 more rows
```

Read a csv file and convert to a tibble

```
prison <- readr::read_csv("data/prison_population.csv") %>%  
  mutate(Quarter = yearquarter(date))
```

```
## # A tibble: 3,072 x 7
```

##	date	state	gender	legal	indigenous	count	Quarter
##	<date>	<chr>	<chr>	<chr>	<chr>	<dbl>	<qtr>
## 1	2005-03-01	ACT	Female	Remanded	ATSI	0	2005 Q1
## 2	2005-03-01	ACT	Female	Remanded	Other	2	2005 Q1
## 3	2005-03-01	ACT	Female	Sentenc~	ATSI	0	2005 Q1
## 4	2005-03-01	ACT	Female	Sentenc~	Other	0	2005 Q1
## 5	2005-03-01	ACT	Male	Remanded	ATSI	7	2005 Q1
## 6	2005-03-01	ACT	Male	Remanded	Other	58	2005 Q1
## 7	2005-03-01	ACT	Male	Sentenc~	ATSI	0	2005 Q1
## 8	2005-03-01	ACT	Male	Sentenc~	Other	0	2005 Q1
## 9	2005-03-01	NSW	Female	Remanded	ATSI	51	2005 Q1
## 10	2005-03-01	NSW	Female	Remanded	Other	131	2005 Q1
## #	... with 3,062 more rows						

Read a csv file and convert to a tibble

```
prison <- readr::read_csv("data/prison_population.csv") %>%  
  mutate(Quarter = yearquarter(date)) %>%  
  select(-date)
```

A tibble: 3,072 x 6

##	state	gender	legal	indigenous	count	Quarter
##	<chr>	<chr>	<chr>	<chr>	<dbl>	<qtr>
## 1	ACT	Female	Remanded	ATSI	0	2005 Q1
## 2	ACT	Female	Remanded	Other	2	2005 Q1
## 3	ACT	Female	Sentenced	ATSI	0	2005 Q1
## 4	ACT	Female	Sentenced	Other	0	2005 Q1
## 5	ACT	Male	Remanded	ATSI	7	2005 Q1
## 6	ACT	Male	Remanded	Other	58	2005 Q1
## 7	ACT	Male	Sentenced	ATSI	0	2005 Q1
## 8	ACT	Male	Sentenced	Other	0	2005 Q1
## 9	NSW	Female	Remanded	ATSI	51	2005 Q1
## 10	NSW	Female	Remanded	Other	131	2005 Q1

with 2,062 more rows

Read a csv file and convert to a tsibble

```
prison <- readr::read_csv("data/prison_population.csv") %>%  
  mutate(Quarter = yearquarter(date)) %>%  
  select(-date) %>%  
  as_tsibble(  
    index = Quarter,  
    key = c(state, gender, legal, indigenous)  
  )
```

```
## # A tsibble: 3,072 x 6 [1Q]  
## # Key:      state, gender, legal, indigenous [64]  
##   state gender legal   indigenous count Quarter  
##   <chr> <chr>  <chr>    <chr>      <dbl>   <qtr>  
## 1 ACT   Female Remanded ATSI         0 2005 Q1  
## 2 ACT   Female Remanded ATSI         1 2005 Q2  
## 3 ACT   Female Remanded ATSI         0 2005 Q3  
## 4 ACT   Female Remanded ATSI         0 2005 Q4  
## 5 ACT   Female Remanded ATSI         1 2006 Q1  
## 6 ACT   Female Remanded ATSI         1 2006 Q2
```


Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

Australian Pharmaceutical Benefits Scheme



Australian Pharmaceutical Benefits Scheme

The **Pharmaceutical Benefits Scheme** (PBS) is the Australian government drugs subsidy scheme.

Australian Pharmaceutical Benefits Scheme

The **Pharmaceutical Benefits Scheme** (PBS) is the Australian government drugs subsidy scheme.

- Many drugs bought from pharmacies are subsidised to allow more equitable access to modern drugs.
- The cost to government is determined by the number and types of drugs purchased. Currently nearly 1% of GDP.
- The total cost is budgeted based on forecasts of drug usage.
- Costs are disaggregated by drug type (ATC1 x15 / ATC2 84), concession category (x2) and patient type (x2), giving $84 \times 2 \times 2 = 336$ time series.

Working with tsibble objects

PBS

```
## # A tsibble: 65,219 x 9 [1M]
## # Key:      Concession, Type, ATC1, ATC2 [336]
##      Month Concession  Type  ATC1  ATC1_desc  ATC2  ATC2_desc  Scripts  Cost
##      <mth> <chr>      <chr> <chr> <chr>      <chr> <chr>      <dbl> <dbl>
##  1 1991 Jul  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 18228 67877
##  2 1991 Aug  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 15327 57011
##  3 1991 Sep  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 14775 55020
##  4 1991 Oct  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 15380 57222
##  5 1991 Nov  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 14371 52120
##  6 1991 Dec  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 15028 54299
##  7 1992 Jan  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 11040 39753
##  8 1992 Feb  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 15165 54405
##  9 1992 Mar  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 16898 61108
## 10 1992 Apr  Concession~ Co-pa~ A      Alimentar~ A01    STOMATOLO~ 18141 65356
## # ... with 65,209 more rows
```

Working with tsibble objects

We can use the `filter()` function to select rows.

```
PBS %>%  
  filter(ATC2 == "A10")
```

```
## # A tsibble: 816 x 9 [1M]  
## # Key:      Concession, Type, ATC1, ATC2 [4]  
##      Month Concession  Type  ATC1  ATC1_desc  ATC2  ATC2_desc  Scripts  Cost  
##      <mt> <chr>      <chr> <chr> <chr>      <chr> <chr>      <dbl> <dbl>  
## 1 1991 Jul  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 89733 2.09e6  
## 2 1991 Aug  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 77101 1.80e6  
## 3 1991 Sep  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 76255 1.78e6  
## 4 1991 Oct  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 78681 1.85e6  
## 5 1991 Nov  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 70554 1.69e6  
## 6 1991 Dec  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 75814 1.84e6  
## 7 1992 Jan  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 64186 1.56e6  
## 8 1992 Feb  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 75899 1.73e6  
## 9 1992 Mar  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 89445 2.05e6  
## 10 1992 Apr  Concession~ Co-pa~ A      Alimentar~ A10    ANTIDIAB~ 97315 2.23e6
```

Working with tsibble objects

We can use the `select()` function to select columns.

```
PBS %>%  
  filter(ATC2 == "A10") %>%  
  select(Month, Concession, Type, Cost)
```

```
## # A tsibble: 816 x 4 [1M]  
## # Key:      Concession, Type [4]  
##      Month Concession  Type          Cost  
##      <mtm> <chr>      <chr>          <dbl>  
##  1 1991 Jul Concessional Co-payments 2092878  
##  2 1991 Aug Concessional Co-payments 1795733  
##  3 1991 Sep Concessional Co-payments 1777231  
##  4 1991 Oct Concessional Co-payments 1848507  
##  5 1991 Nov Concessional Co-payments 1686458  
##  6 1991 Dec Concessional Co-payments 1843079  
##  7 1992 Jan Concessional Co-payments 1564702  
##  8 1992 Feb Concessional Co-payments 1732508  
##  9 1992 Mar Concessional Co-payments 2046102  
## 10 1992 Apr Concessional Co-payments 2225077
```

Working with `tsibble` objects

We can use the `summarise()` function to summarise over keys.

```
PBS %>%  
  filter(ATC2 == "A10") %>%  
  select(Month, Concession, Type, Cost) %>%  
  summarise(total_cost = sum(Cost))
```

```
## # A tsibble: 204 x 2 [1M]
```

```
##       Month total_cost
```

```
##       <mth>       <dbl>
```

```
## 1 1991 Jul      3526591
```

```
## 2 1991 Aug      3180891
```

```
## 3 1991 Sep      3252221
```

```
## 4 1991 Oct      3611003
```

```
## 5 1991 Nov      3565869
```

```
## 6 1991 Dec      4306371
```

```
## 7 1992 Jan      5088335
```

```
## 8 1992 Feb      2814520
```

```
## 9 1992 Mar      2985811
```

```
## 10 1992 Apr      2204700
```


Working with tsibble objects

We can use the `mutate()` function to create new variables.

```
PBS %>%  
  filter(ATC2 == "A10") %>%  
  select(Month, Concession, Type, Cost) %>%  
  summarise(total_cost = sum(Cost)) %>%  
  mutate(total_cost = total_cost / 1e6)
```

```
## # A tsibble: 204 x 2 [1M]
```

```
##       Month total_cost
```

```
##       <mth>      <dbl>
```

```
## 1 1991 Jul       3.53
```

```
## 2 1991 Aug       3.18
```

```
## 3 1991 Sep       3.25
```

```
## 4 1991 Oct       3.61
```

```
## 5 1991 Nov       3.57
```

```
## 6 1991 Dec       4.31
```

```
## 7 1992 Jan       5.09
```

```
## 8 1992 Feb       2.81
```

```
## 9 1992 Mar       2.00
```

Working with tsibble objects

We can use the `mutate()` function to create new variables.

```
PBS %>%  
  filter(ATC2 == "A10") %>%  
  select(Month, Concession, Type, Cost) %>%  
  summarise(total_cost = sum(Cost)) %>%  
  mutate(total_cost = total_cost / 1e6) -> a10
```

```
## # A tsibble: 204 x 2 [1M]
```

```
##       Month total_cost
```

```
##       <mtch>      <dbl>
```

```
##  1 1991 Jul         3.53
```

```
##  2 1991 Aug         3.18
```

```
##  3 1991 Sep         3.25
```

```
##  4 1991 Oct         3.61
```

```
##  5 1991 Nov         3.57
```

```
##  6 1991 Dec         4.31
```

```
##  7 1992 Jan         5.09
```

```
##  8 1992 Feb         2.81
```

```
##  9 1992 Mar         2.00
```

Your turn

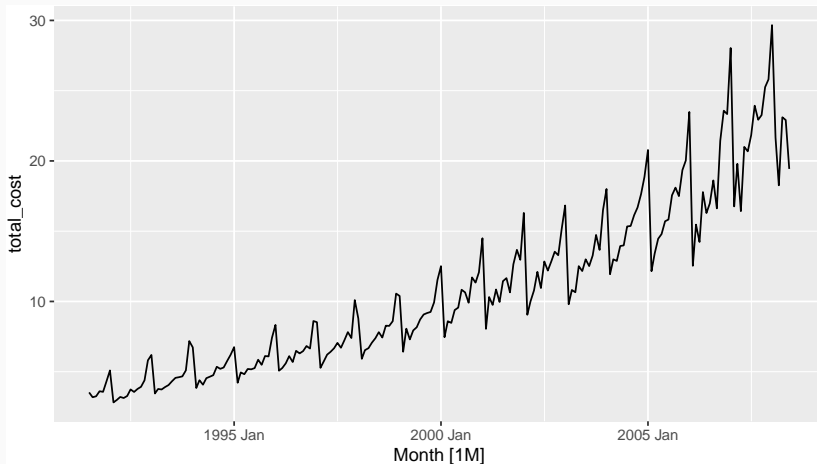
- 1 Download `tourism.xlsx` from <http://robjhyndman.com/data/tourism.xlsx>, and read it into R using `read_excel()` from the `readxl` package.
- 2 Create a `tsibble` which is identical to the `tourism` `tsibble` from the `tsibble` package.
- 3 Find what combination of `Region` and `Purpose` had the maximum number of overnight trips on average.
- 4 Create a new `tsibble` which combines the `Purposes` and `Regions`, and just has total trips by `State`.

Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

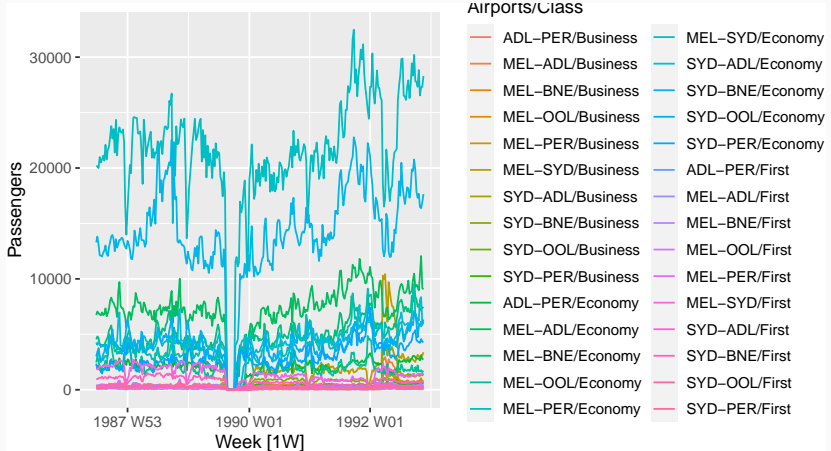
Time plots

```
a10 %>%  
  autoplot(total_cost)
```



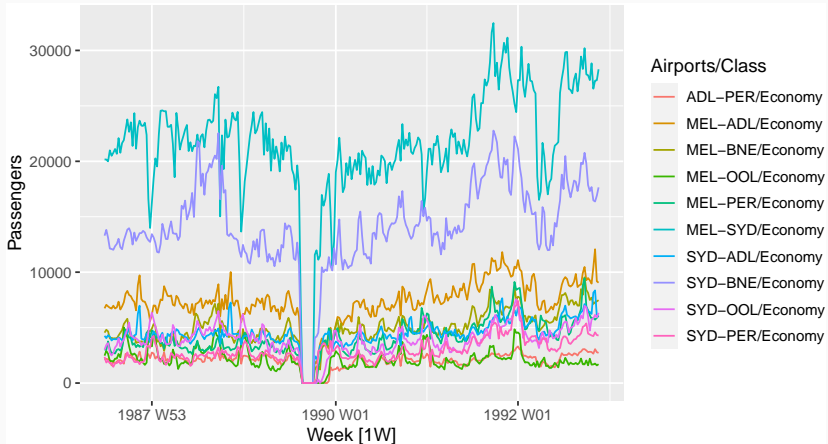
Ansett airlines

```
ansett %>%
  autoplot(Passengers)
```



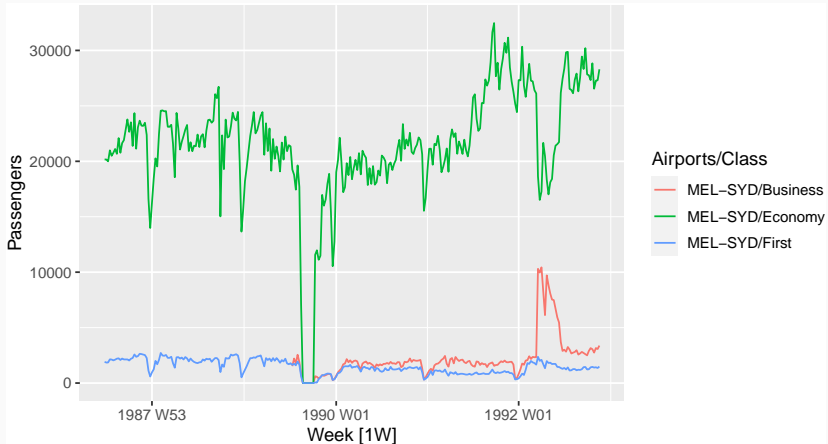
Ansett airlines

```
ansett %>%  
  filter(Class == "Economy") %>%  
  autoplot(Passengers)
```



Ansett airlines

```
ansett %>%  
  filter(Airports == "MEL-SYD") %>%  
  autoplot(Passengers)
```



Your turn

- Create plots of the following time series: Bricks from `aus_production`, Lynx from `pel_t`, Close from `gafa_stock`, Demand from `vic_elec`.
- Use `help()` to find out about the data in each series.
- For the last plot, modify the axis labels and title.

Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

Time series patterns

Trend pattern exists when there is a long-term increase or decrease in the data.

Seasonal pattern exists when a series is influenced by seasonal factors (e.g., the quarter of the year, the month, or day of the week).

Cyclic pattern exists when data exhibit rises and falls that are *not of fixed period* (duration usually of at least 2 years).

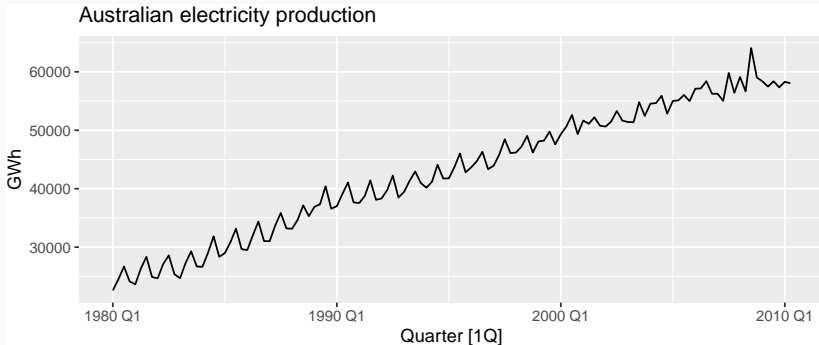
Time series components

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

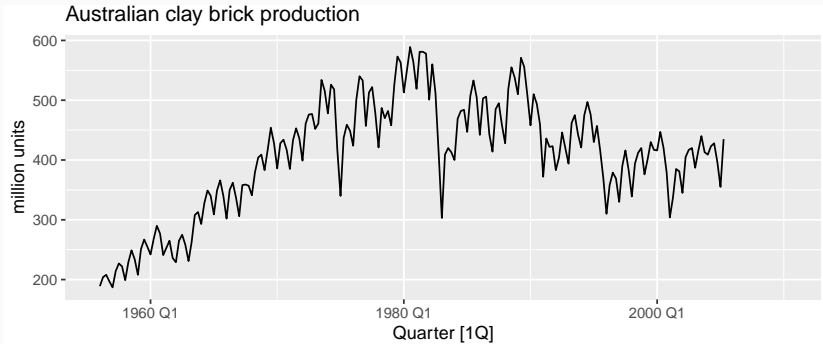
Time series patterns

```
aus_production %>%  
  filter(year(Quarter) >= 1980) %>%  
  autoplot(Electricity) +  
  labs(y = "GWh",  
       title = "Australian electricity production")
```



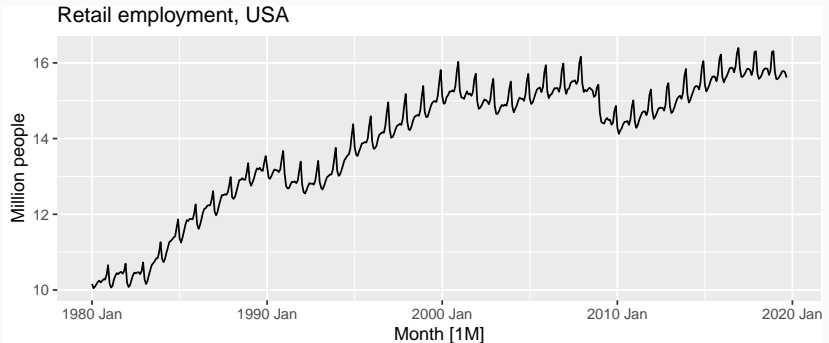
Time series patterns

```
aus_production %>%  
  autoplot(Bricks) +  
  labs(y = "million units",  
       title = "Australian clay brick production")
```



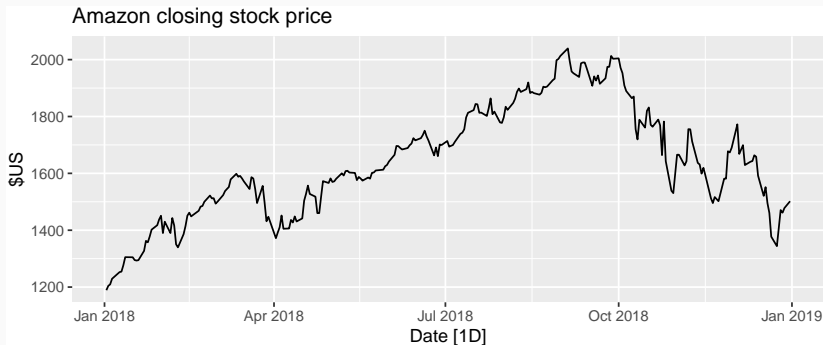
Time series patterns

```
us_employment %>%  
  filter(Title == "Retail Trade", year(Month) >= 1980) %>%  
  autoplot(Employed / 1e3) +  
  labs(y = "Million people",  
       title = "Retail employment, USA")
```



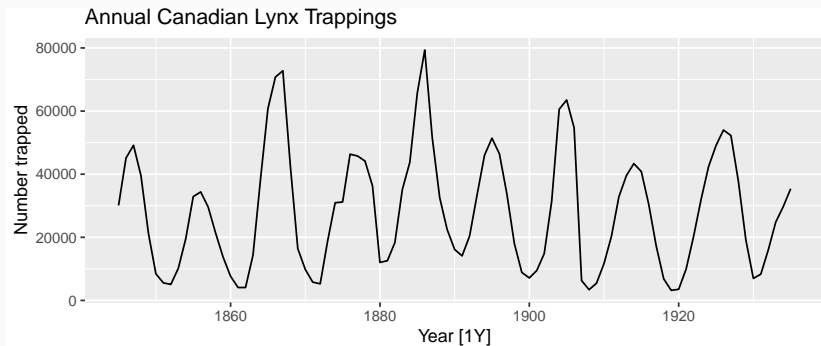
Time series patterns

```
gafa_stock %>%  
  filter(Symbol == "AMZN", year(Date) >= 2018) %>%  
  autoplot(Close) +  
  labs(y = "$US",  
       title = "Amazon closing stock price")
```



Time series patterns

```
pelt %>%  
  autoplot(Lynx) +  
  labs(y="Number trapped",  
       title = "Annual Canadian Lynx Trappings")
```



Seasonal or cyclic?

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

Seasonal or cyclic?

Differences between seasonal and cyclic patterns:

- seasonal pattern constant length; cyclic pattern variable length
- average length of cycle longer than length of seasonal pattern
- magnitude of cycle more variable than magnitude of seasonal pattern

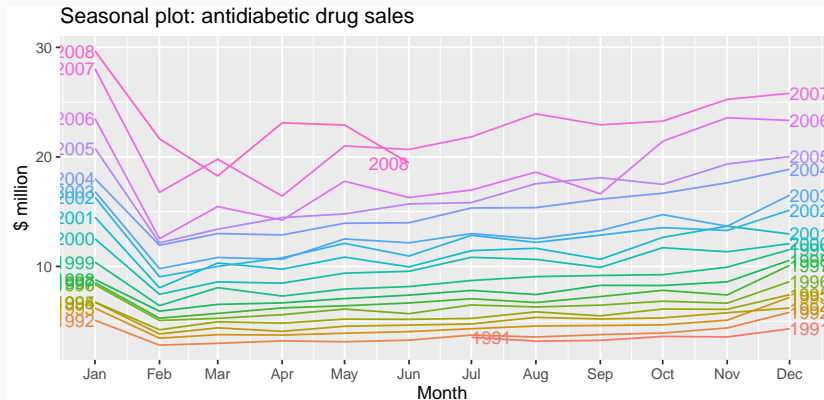
The timing of peaks and troughs is predictable with seasonal data, but unpredictable in the long term with cyclic data.

Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

Seasonal plots

```
a10 %>% gg_season(total_cost, labels = "both") +  
  labs(y = "$ million",  
       title = "Seasonal plot: antidiabetic drug sales")
```

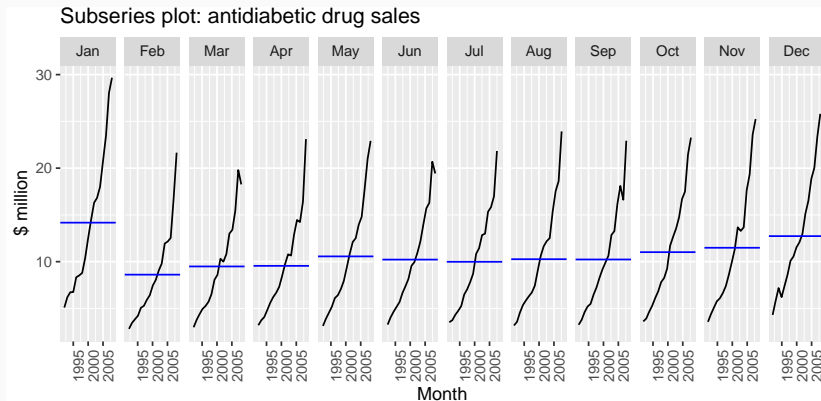


Seasonal plots

- Data plotted against the individual “seasons” in which the data were observed. (In this case a “season” is a month.)
- Something like a time plot except that the data from each season are overlapped.
- Enables the underlying seasonal pattern to be seen more clearly, and also allows any substantial departures from the seasonal pattern to be easily identified.
- In R: `gg_season()`

Seasonal subseries plots

```
a10 %>%  
  gg_subseries(total_cost) +  
  labs(y = "$ million",  
       title = "Subseries plot: antidiabetic drug sales")
```

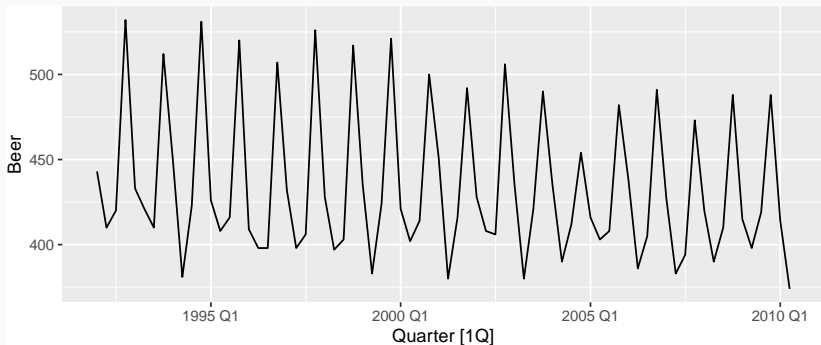


Seasonal subseries plots

- Data for each season collected together in time plot as separate time series.
- Enables the underlying seasonal pattern to be seen clearly, and changes in seasonality over time to be visualized.
- In R: `gg_subseries()`

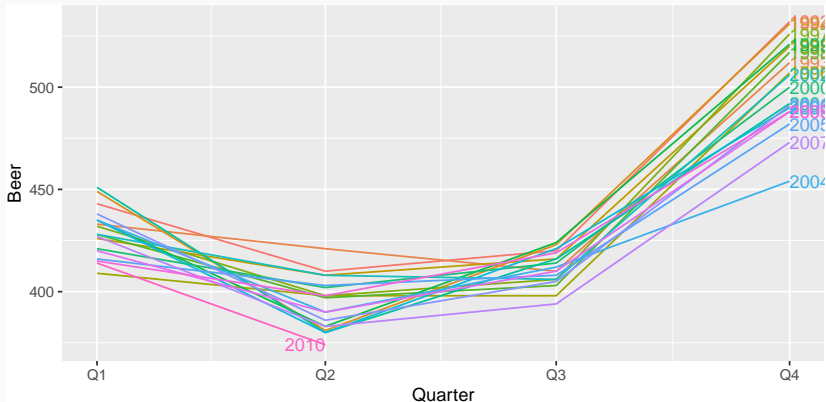
Quarterly Australian Beer Production

```
beer <- aus_production %>%  
  select(Quarter, Beer) %>%  
  filter(year(Quarter) >= 1992)  
beer %>% autoplot(Beer)
```



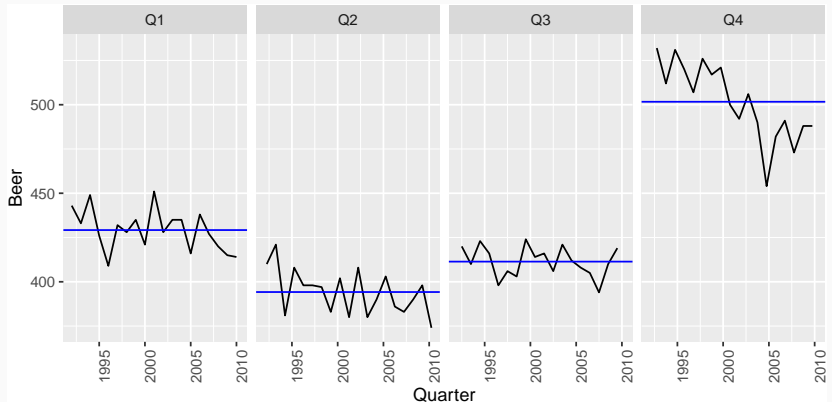
Quarterly Australian Beer Production

```
beer %>% gg_season(Beer, labels="right")
```



Quarterly Australian Beer Production

```
beer %>% gg_subseries(Beer)
```



Your turn

Look at the quarterly tourism data for the Snowy Mountains

```
snowy <- tourism %>%  
  filter(Region == "Snowy Mountains")
```

- Use `autoplot()`, `gg_season()` and `gg_subseries()` to explore the data.
- What do you learn?

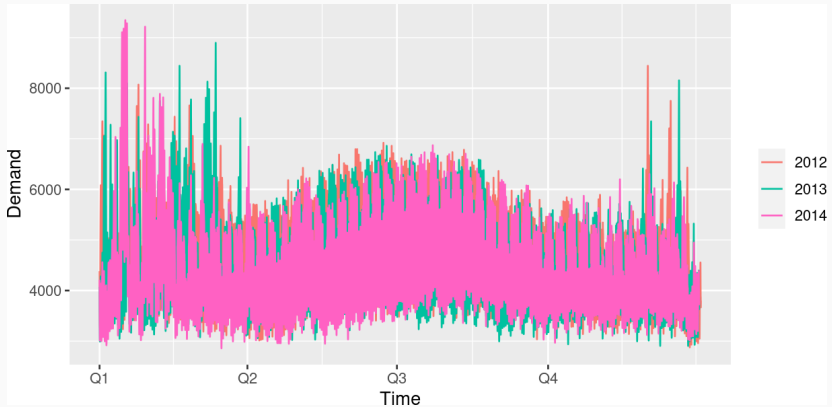
Multiple seasonal periods

```
vic_elec
```

```
## # A tibble: 52,608 x 5 [30m] <Australia/Melbourne>
##   Time                Demand Temperature Date        Holiday
##   <dtm>                <dbl>         <dbl> <date>        <lgl>
## 1 2012-01-01 00:00:00  4383.          21.4 2012-01-01 TRUE
## 2 2012-01-01 00:30:00  4263.          21.0 2012-01-01 TRUE
## 3 2012-01-01 01:00:00  4049.          20.7 2012-01-01 TRUE
## 4 2012-01-01 01:30:00  3878.          20.6 2012-01-01 TRUE
## 5 2012-01-01 02:00:00  4036.          20.4 2012-01-01 TRUE
## 6 2012-01-01 02:30:00  3866.          20.2 2012-01-01 TRUE
## 7 2012-01-01 03:00:00  3694.          20.1 2012-01-01 TRUE
## 8 2012-01-01 03:30:00  3562.          19.6 2012-01-01 TRUE
## 9 2012-01-01 04:00:00  3433.          19.1 2012-01-01 TRUE
## 10 2012-01-01 04:30:00  3359.          19.0 2012-01-01 TRUE
## # ... with 52,598 more rows
```

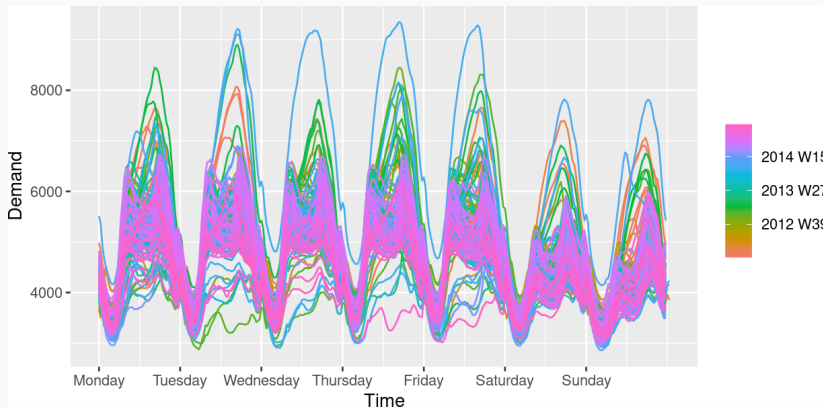
Multiple seasonal periods

```
vic_elec %>% gg_season(Demand)
```



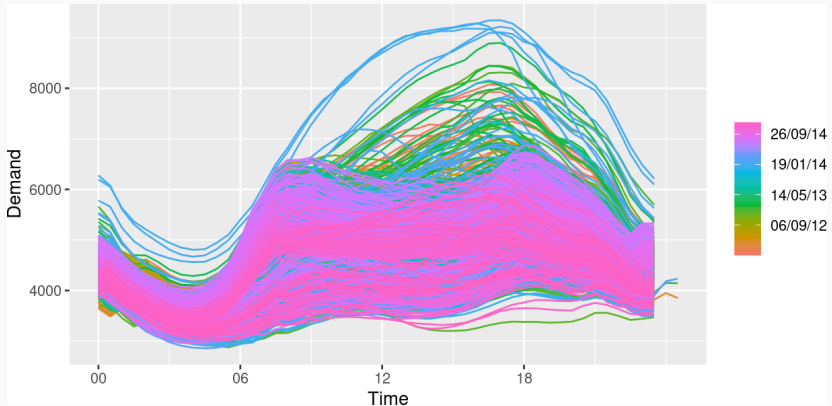
Multiple seasonal periods

```
vic_elec %>% gg_season(Demand, period = "week")
```



Multiple seasonal periods

```
vic_elec %>% gg_season(Demand, period = "day")
```



Australian holidays

```
holidays <- tourism %>%  
  filter(Purpose == "Holiday") %>%  
  group_by(State) %>%  
  summarise(Trips = sum(Trips))
```

```
## # A tsibble: 640 x 3 [1Q]
```

```
## # Key:           State [8]
```

```
##   State Quarter Trips
```

```
##   <chr>    <qtr> <dbl>
```

```
##  1 ACT     1998 Q1  196.
```

```
##  2 ACT     1998 Q2  127.
```

```
##  3 ACT     1998 Q3  111.
```

```
##  4 ACT     1998 Q4  170.
```

```
##  5 ACT     1999 Q1  108.
```

```
##  6 ACT     1999 Q2  125.
```

```
##  7 ACT     1999 Q3  178.
```

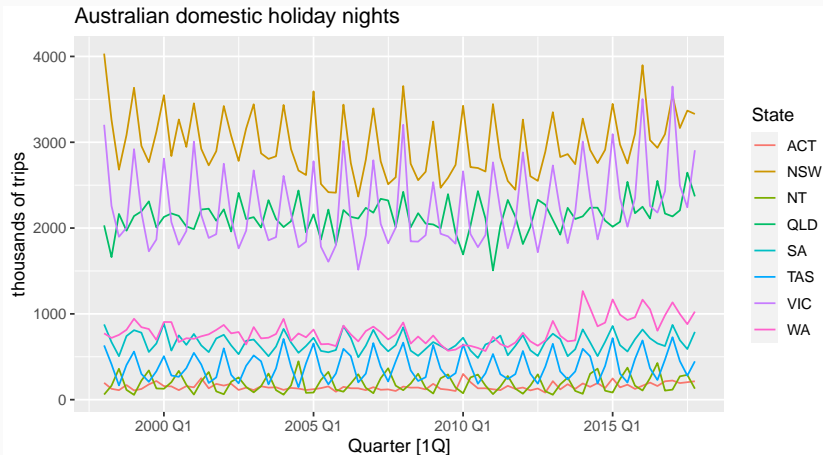
```
##  8 ACT     1999 Q4  218.
```

```
##  9 ACT     2000 Q1  158.
```

```
## 10 ACT     2000 Q2  155.
```

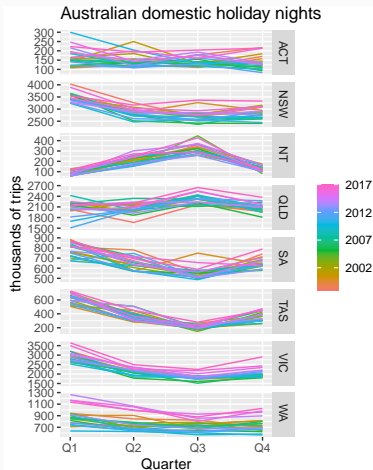
Australian holidays

```
holidays %>% autoplot(Trips) +  
  labs(y = "thousands of trips",  
       title = "Australian domestic holiday nights")
```



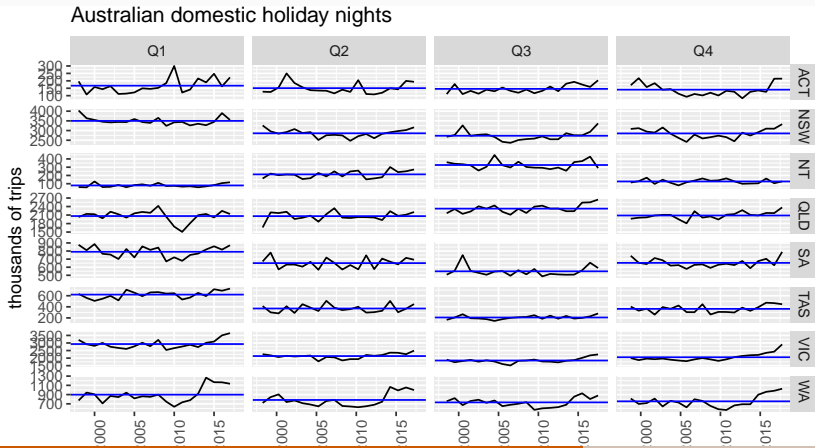
Seasonal plots

```
holidays %>% gg_season(Trips) +  
  labs(y = "thousands of trips",  
       title = "Australian domestic holiday nights")
```



Seasonal subseries plots

```
holidays %>%  
  gg_subseries(Trips) +  
    labs(y = "thousands of trips",  
         title = "Australian domestic holiday nights")
```



Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

Example: Beer production

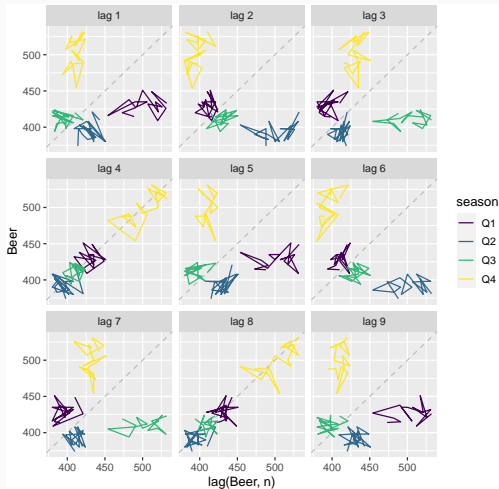
```
new_production <- aus_production %>%  
  filter(year(Quarter) >= 1992)  
new_production
```

```
## # A tsibble: 74 x 7 [1Q]
```

##		Quarter	Beer	Tobacco	Bricks	Cement	Electricity	Gas
##		<qtr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
##	1	1992 Q1	443	5777	383	1289	38332	117
##	2	1992 Q2	410	5853	404	1501	39774	151
##	3	1992 Q3	420	6416	446	1539	42246	175
##	4	1992 Q4	532	5825	420	1568	38498	129
##	5	1993 Q1	433	5724	394	1450	39460	116
##	6	1993 Q2	421	6036	462	1668	41356	149
##	7	1993 Q3	410	6570	475	1648	42949	163
##	8	1993 Q4	512	5675	443	1863	40974	138

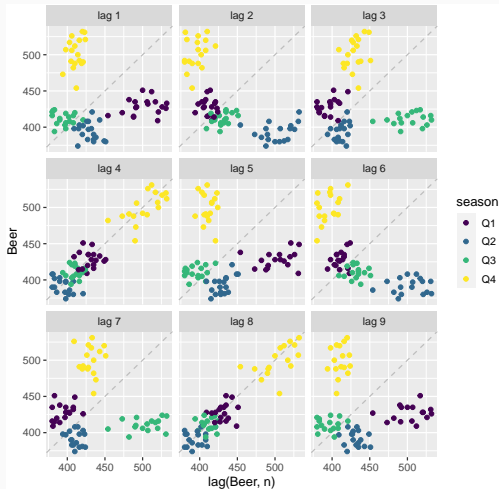
Example: Beer production

```
new_production %>% gg_lag(Beer)
```



Example: Beer production

```
new_production %>% gg_lag(Beer, geom='point')
```



Lagged scatterplots

- Each graph shows y_t plotted against y_{t-k} for different values of k .
- The autocorrelations are the correlations associated with these scatterplots.
- ACF (autocorrelation function):
 - ▶ $r_1 = \text{Correlation}(y_t, y_{t-1})$
 - ▶ $r_2 = \text{Correlation}(y_t, y_{t-2})$
 - ▶ $r_3 = \text{Correlation}(y_t, y_{t-3})$
 - ▶ etc.

Autocorrelation

Covariance and **correlation**: measure extent of **linear relationship** between two variables (y and X).

Autocorrelation

Covariance and **correlation**: measure extent of **linear relationship** between two variables (y and X).

Autocovariance and **autocorrelation**: measure linear relationship between **lagged values** of a time series y .

Autocorrelation

Covariance and **correlation**: measure extent of **linear relationship** between two variables (y and X).

Autocovariance and **autocorrelation**: measure linear relationship between **lagged values** of a time series y .

We measure the relationship between:

- y_t and y_{t-1}
- y_t and y_{t-2}
- y_t and y_{t-3}
- etc.

Autocorrelation

We denote the sample autocovariance at lag k by c_k and the sample autocorrelation at lag k by r_k . Then define

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and $r_k = c_k / c_0$

Autocorrelation

We denote the sample autocovariance at lag k by c_k and the sample autocorrelation at lag k by r_k . Then define

$$c_k = \frac{1}{T} \sum_{t=k+1}^T (y_t - \bar{y})(y_{t-k} - \bar{y})$$

and $r_k = c_k / c_0$

- r_1 indicates how successive values of y relate to each other
- r_2 indicates how y values two periods apart relate to each other
- r_k is *almost* the same as the sample correlation between y_t and y_{t-k} .

Autocorrelation

Results for first 9 lags for beer data:

```
new_production %>% ACF(Beer, lag_max = 9)
```

```
## # A tsibble: 9 x 2 [1Q]
```

```
##   lag    acf
```

```
##   <lag>  <dbl>
```

```
## 1    1Q -0.102
```

```
## 2    2Q -0.657
```

```
## 3    3Q -0.0603
```

```
## 4    4Q  0.869
```

```
## 5    5Q -0.0892
```

```
## 6    6Q -0.635
```

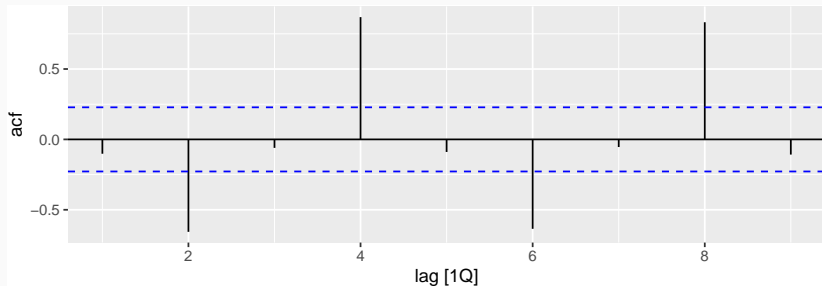
```
## 7    7Q -0.0542
```

```
## 8    8Q  0.832
```

Autocorrelation

Results for first 9 lags for beer data:

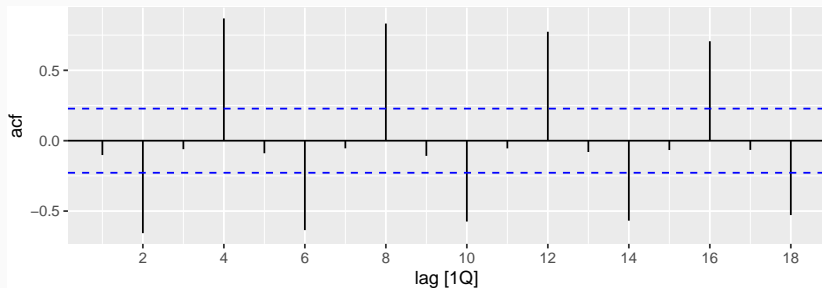
```
new_production %>% ACF(Beer, lag_max = 9) %>% autoplot()
```



- Together, the autocorrelations at lags 1, 2, ..., make up the *autocorrelation* or ACF.
- The plot is known as a **correlogram**

Autocorrelation

```
new_production %>% ACF(Beer) %>% autoplot()
```



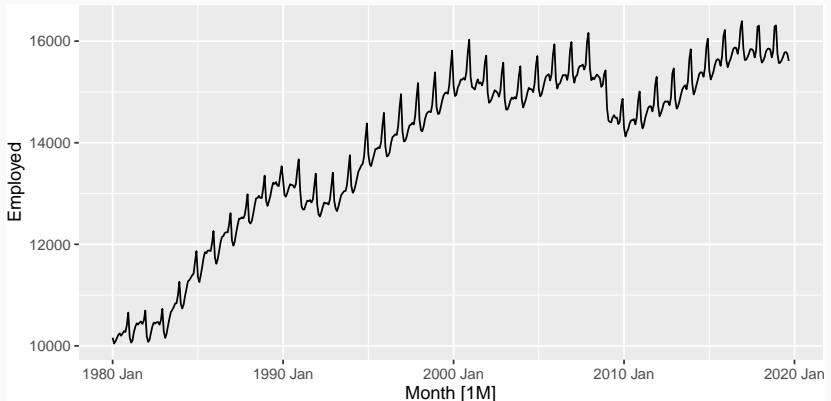
- r_4 higher than for the other lags. This is due to **the seasonal pattern in the data**: the peaks tend to be **4 quarters** apart and the troughs tend to be **2 quarters** apart.
- r_2 is more negative than for the other lags because troughs tend to be 2 quarters behind peaks.

Trend and seasonality in ACF plots

- When data have a trend, the autocorrelations for small lags tend to be large and positive.
- When data are seasonal, the autocorrelations will be larger at the seasonal lags (i.e., at multiples of the seasonal frequency)
- When data are trended and seasonal, you see a combination of these effects.

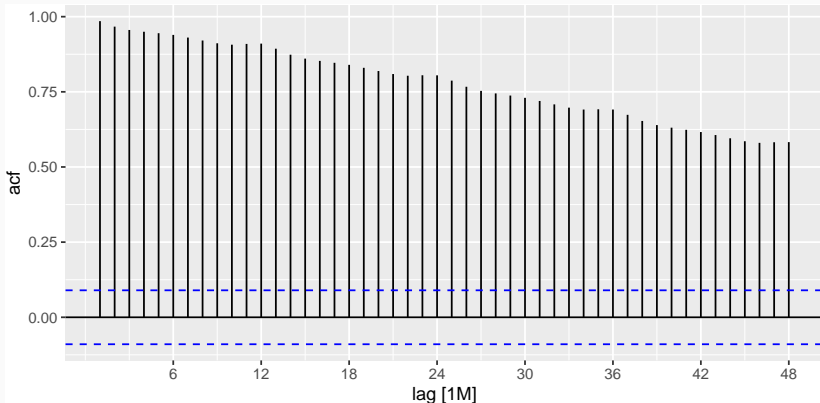
US retail trade employment

```
retail <- us_employment %>%  
  filter(Title == "Retail Trade", year(Month) >= 1980)  
retail %>% autoplot(Employed)
```



US retail trade employment

```
retail %>%  
  ACF(Employed, lag_max = 48) %>%  
  autoplot()
```



Google stock price

```
google_2015 <- gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) == 2015) %>%  
  select(Date, Close)  
google_2015
```

```
## # A tsibble: 252 x 2 [1D]
```

```
##   Date      Close
```

```
##   <date>    <dbl>
```

```
## 1 2015-01-02  522.
```

```
## 2 2015-01-05  511.
```

```
## 3 2015-01-06  499.
```

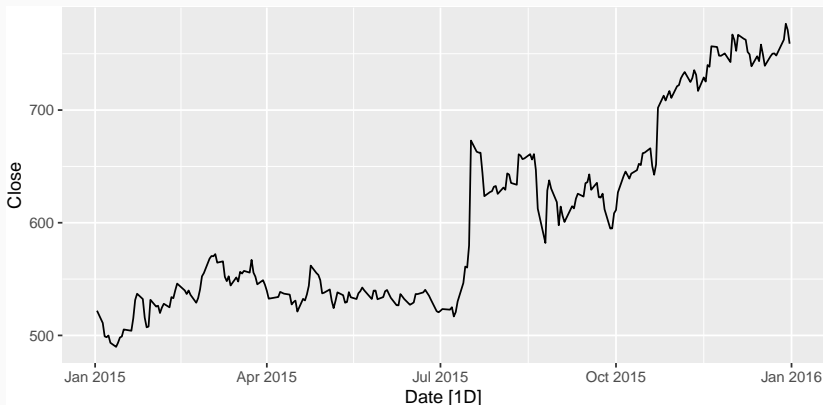
```
## 4 2015-01-07  498.
```

```
## 5 2015-01-08  500.
```

```
## 6 2015-01-09  493.
```

Google stock price

```
google_2015 %>% autoplot(Close)
```



Google stock price

```
google_2015 %>%  
  ACF(Close, lag_max=100)  
# Error: Can't handle tsibble of irregular interval.
```

Google stock price

```
google_2015 %>%  
  ACF(Close, lag_max=100)  
# Error: Can't handle tsibble of irregular interval.
```

```
google_2015
```

```
## # A tsibble: 252 x 2 [1D]
```

```
##   Date      Close
```

```
##   <date>    <dbl>
```

```
## 1 2015-01-02  522.
```

```
## 2 2015-01-05  511.
```

```
## 3 2015-01-06  499.
```

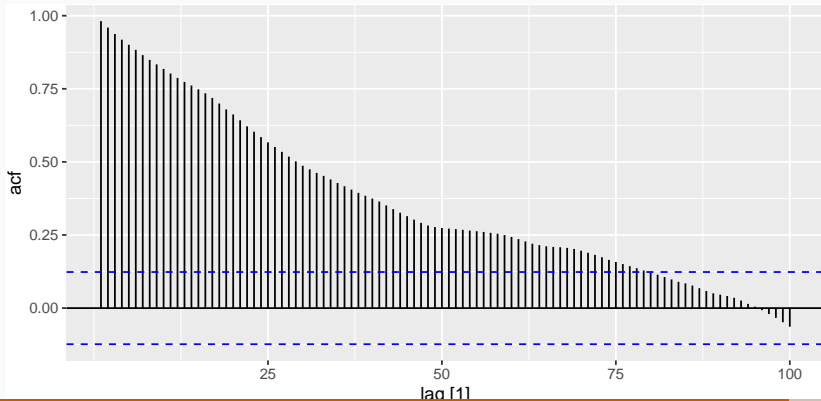

Google stock price

```
google_2015 <- google_2015 %>%  
  mutate(trading_day = row_number()) %>%  
  update_tsibble(index = trading_day, regular = TRUE)  
google_2015
```

```
## # A tsibble: 252 x 3 [1]  
##   Date      Close trading_day  
##   <date>    <dbl>      <int>  
## 1 2015-01-02  522.         1  
## 2 2015-01-05  511.         2  
## 3 2015-01-06  499.         3  
## 4 2015-01-07  498.         4  
## 5 2015-01-08  500.         5  
## 6 2015-01-09  493.         6  
## 7 2015-01-12  490.         7
```

Google stock price

```
google_2015 %>%  
  ACF(Close, lag_max = 100) %>%  
  autoplot()
```



Your turn

We have introduced the following functions:

- `gg_lag`
- `ACF`

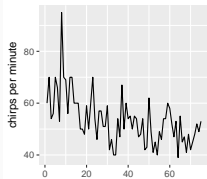
Use these functions to explore the following time series:

- Bricks from `aus_production`
- Lynx from `pelt`
- Victorian Electricity Demand from `aus_elec`

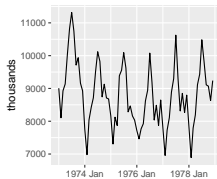
Can you spot any seasonality, cyclicity and trend?
What do you learn about the series?

Which is which?

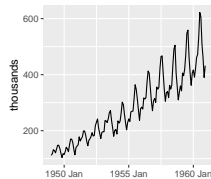
1. Daily temperature of cow



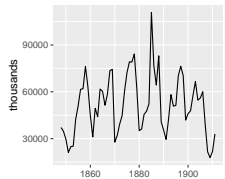
2. Monthly accidental deaths



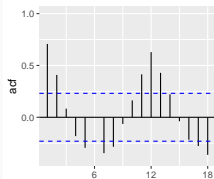
3. Monthly air passengers



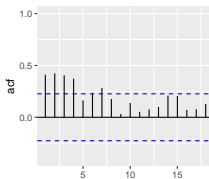
4. Annual mink trappings



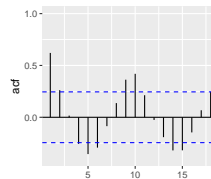
A



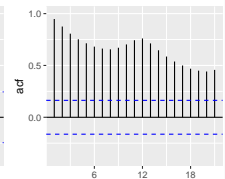
B



C



D

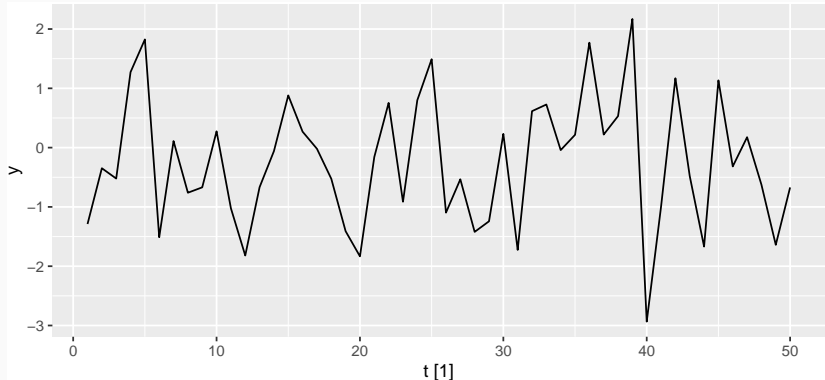


Outline

- 1 Time series in R
- 2 Example: Australian prison population
- 3 Example: Australian pharmaceutical sales
- 4 Time plots
- 5 Time series patterns
- 6 Seasonal and subseries plots
- 7 Lag plots and autocorrelation
- 8 White noise

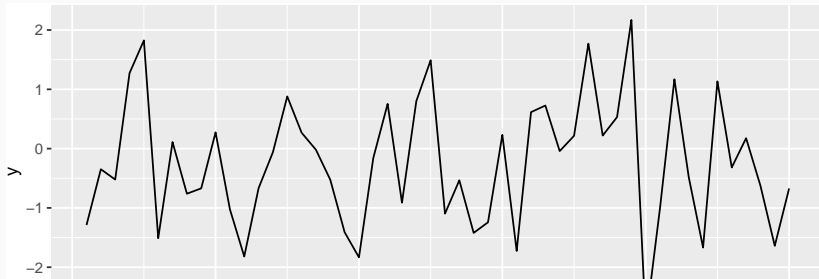
Example: White noise

```
set.seed(30)  
wn <- tsibble(t = 1:50, y = rnorm(50), index = t)  
wn %>% autoplot(y)
```



Example: White noise

```
set.seed(30)  
wn <- tsibble(t = 1:50, y = rnorm(50), index = t)  
wn %>% autoplot(y)
```

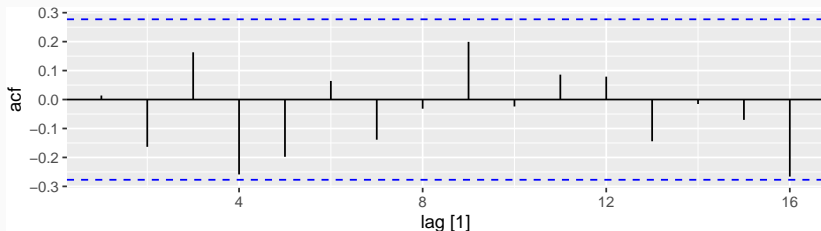


White noise data is uncorrelated across time with zero mean and constant variance.
(Technically, we require independence as well.)

Example: White noise

wn %>% ACF(y)

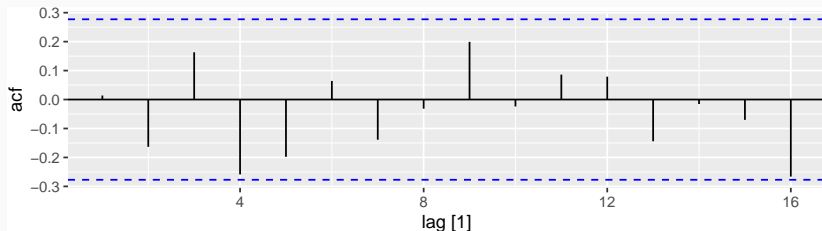
r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9	r_{10}
0.014	-0.163	0.163	-0.259	-0.198	0.064	-0.139	-0.032	0.199	-0.024



Example: White noise

wn %>% ACF(y)

r_1	r_2	r_3	r_4	r_5	r_6	r_7	r_8	r_9	r_{10}
0.014	-0.163	0.163	-0.259	-0.198	0.064	-0.139	-0.032	0.199	-0.024



- Sample autocorrelations for white noise series.
- Expect each autocorrelation to be close to zero.
- Blue lines show 95% critical values.

Sampling distribution of autocorrelations

Sampling distribution of r_k for white noise data is asymptotically $N(0, 1/T)$.

Sampling distribution of autocorrelations

Sampling distribution of r_k for white noise data is asymptotically $N(0, 1/T)$.

- 95% of all r_k for white noise must lie within $\pm 1.96/\sqrt{T}$.
- If this is not the case, the series is probably not WN.
- Common to plot lines at $\pm 1.96/\sqrt{T}$ when plotting ACF. These are the **critical values**.

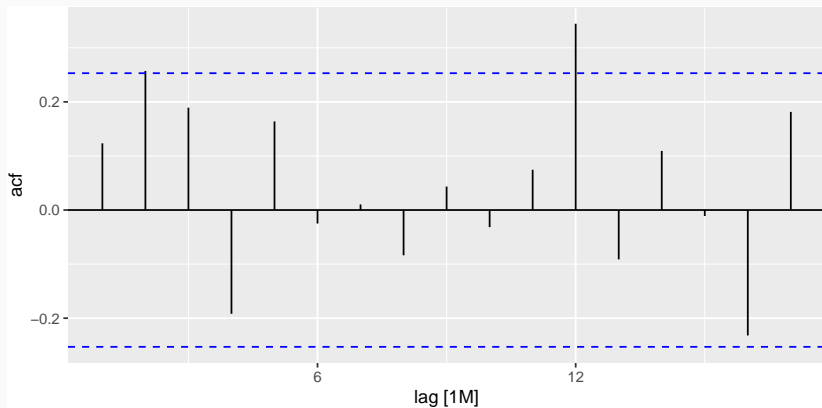
Example: Pigs slaughtered

```
pigs <- aus_livestock %>%  
  filter(State == "Victoria", Animal == "Pigs",  
         year(Month) >= 2014)  
pigs %>% autoplot(Count/1e3) +  
  labs(y = "Thousands",  
       title = "Number of pigs slaughtered in Victoria")
```



Example: Pigs slaughtered

```
pigs %>% ACF(Count) %>% autoplot()
```



Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 2014 through December 2018 (Source: Australian Bureau of Statistics.)

Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 2014 through December 2018 (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows significant autocorrelation for lag 2 and 12.
- Indicate some slight seasonality.

Example: Pigs slaughtered

Monthly total number of pigs slaughtered in the state of Victoria, Australia, from January 2014 through December 2018 (Source: Australian Bureau of Statistics.)

- Difficult to detect pattern in time plot.
- ACF shows significant autocorrelation for lag 2 and 12.
- Indicate some slight seasonality.

These show the series is **not a white noise series**.

Your turn

You can compute the daily changes in the Google stock price in 2018 using

```
dgoog <- gafa_stock %>%  
  filter(Symbol == "GOOG", year(Date) >= 2018) %>%  
  mutate(trading_day = row_number()) %>%  
  update_tsibble(index=trading_day, regular=TRUE) %>%  
  mutate(diff = difference(Close))
```

Does `diff` look like white noise?